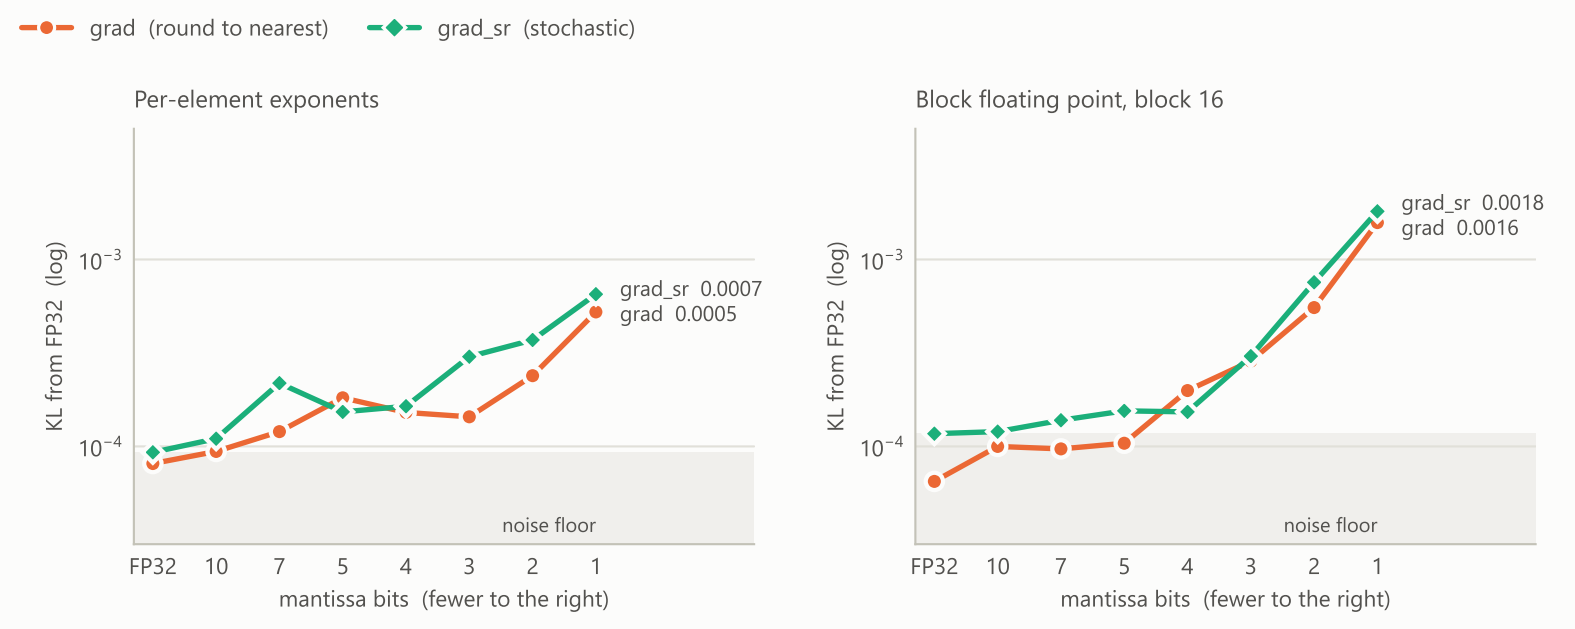
grad (round to nearest)
grad_sr (stochastic)
Per-element exponents
KL from FP32 (log)
10^{-3}
10^{-4}
grad_sr 0.0007
grad 0.0005
noise floor
FP32 10 7 5 4 3 2 1
mantissa bits (fewer to the right)
Block floating point, block 16
KL from FP32 (log)
10^{-3}
10^{-4}
grad_sr 0.0018
grad 0.0016
noise floor
FP32 10 7 5 4 3 2 1
mantissa bits (fewer to the right)