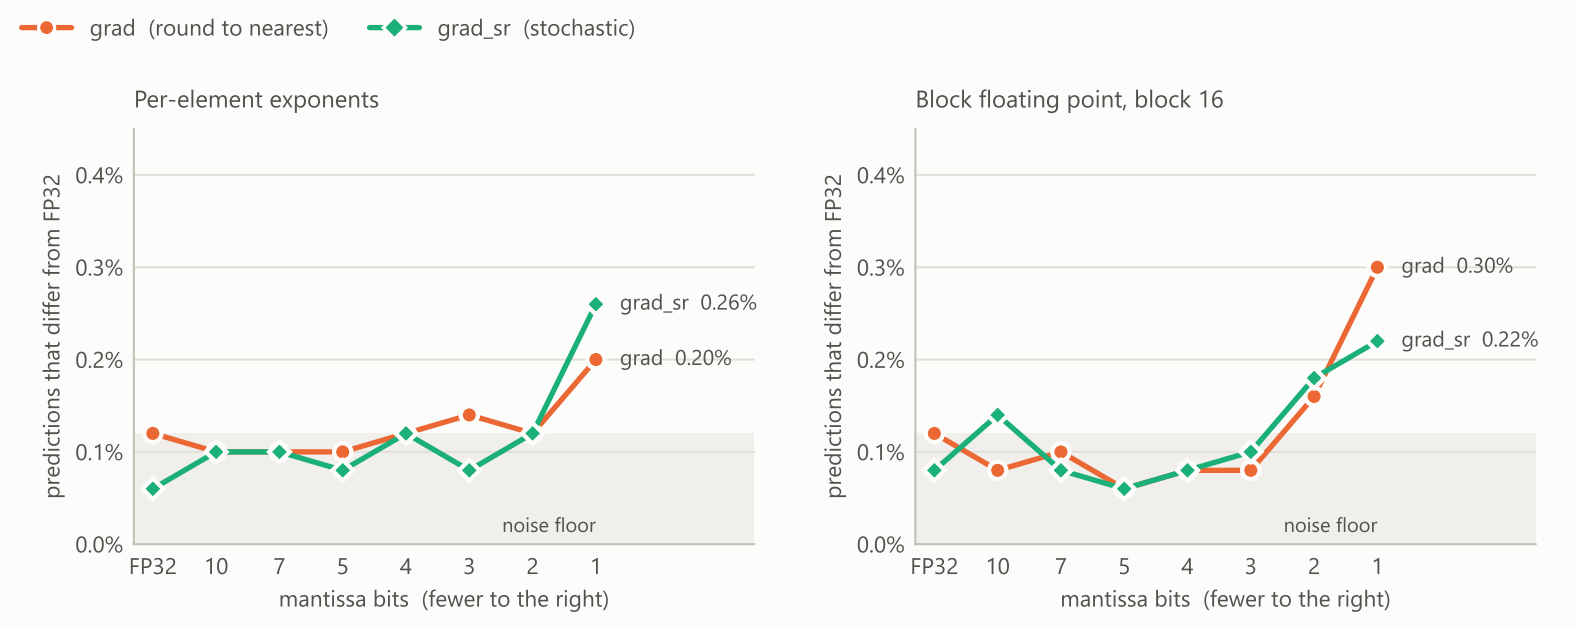
grad (round to nearest)
grad_sr (stochastic)
Per-element exponents
predictions that differ from FP32
0.4%
0.3%
0.2%
0.1%
0.0%
grad_sr 0.26%
grad 0.20%
noise floor
FP32
10
7
5
4
3
2
1
mantissa bits (fewer to the right)
Block floating point, block 16
predictions that differ from FP32
0.4%
0.3%
0.2%
0.1%
0.0%
grad 0.30%
grad_sr 0.22%
noise floor
FP32
10
7
5
4
3
2
1
mantissa bits (fewer to the right)